

Linear Regression Algorithms Demo



LAST UPDATED	LANGUAGE
21 Dec 2017, 2:30 AM	Python 2.7

This notebook shows:

- Install SystemML Python package and jar file
 - pip
 - SystemML 'Hello World'
- Example 1: Matrix Multiplication
- Load diabetes dataset from scikit-learn
- Example 2: Implement three different algorithms to train linear regression model
 - Algorithm 1: Linear Regression - Direct Solve (no regularization)
 - Algorithm 2: Linear Regression - Batch Gradient Descent (no regularization)
 - Algorithm 3: Linear Regression - Conjugate Gradient (no regularization)
- Example 3: Invoke existing SystemML algorithm script LinearRegDS.dml using MLContext API
- Example 4: Invoke existing SystemML algorithm using scikit-learn/SparkML pipeline like API
- Example 5: Invoking a Keras model with SystemML

Install SystemML Python package and jar file

```
In [ ]: #!pip install --upgrade systemml  
!pip install --upgrade https://github.com/niketanpansare/future_of_data/raw/master/systemml-1.1.0-SNAPSHOT-python.tar.gz  
!ln -s -f ~/.local/lib/python2.7/site-packages/systemml/systemml-java/*.jar ~/data/libs/
```

```
In [1]: !pip show systemml  
  
Name: systemml  
Version: 1.1.0  
Summary: Apache SystemML is a distributed and declarative machine learning platform.  
Home-page: http://systemml.apache.org/  
Author: Apache SystemML  
Author-email: dev@systemml.apache.org  
License: Apache 2.0  
Location: /gpfs/global_fs01/sym_shared/YPPProdSpark/user/scf4-b69284e1625908-5ca7710237a9/.local/lib/python2.7/site-packages  
Requires: Pillow, scikit-learn, pandas, scipy, numpy
```

Import SystemML API

```
In [2]: sc.version
```

```
Out[2]: u'2.1.0'
```

```
In [3]: from systemml import MLContext, dml  
# Create a MLContext object  
ml = MLContext(sc)  
# And print the information of SystemML version  
print(ml.info())
```

Archiver-Version: Plexus Archiver

Artifact-Id: systemml
Build-Jdk: 1.8.0_111
Build-Time: 2017-12-19 13:17:52 CST
Built-By: biuser
Created-By: Apache Maven 3.0.5
Group-Id: org.apache.systemml
Main-Class: org.apache.sysml.api.DMLScript
Manifest-Version: 1.0
Minimum-Recommended-Spark-Version: 2.1.0
Version: 1.1.0-SNAPSHOT

```
In [4]: # Create a DML script for a Hello World' example and execute it using MLContext
script = dml("""
print('Hello World');
""")
ml.execute(script)
```

Hello World
SystemML Statistics:
Total execution time: 0.001 sec.
Number of executed Spark inst: 0.

Out[4]: MLResults

```
In [5]: # Let's modify the above script to get the Hello World string
script = dml("""
s = 'Hello World'
""").output("s")

hello_world_str = ml.execute(script).get("s")

print(hello_world_str)
```

SystemML Statistics:
Total execution time: 0.000 sec.
Number of executed Spark inst: 0.

Hello World

Import numpy, sklearn, and define some helper functions

```
In [6]: import sys, os
import matplotlib.pyplot as plt
import numpy as np
from sklearn import datasets
plt.switch_backend('agg')
```

Example 1: Matrix Multiplication

SystemML script to generate a random matrix, perform matrix

SystemML script to generate a random matrix, perform matrix multiplication, and compute the sum of the output

```
In [7]: script = """
        X = rand(rows=$nr, cols=1000, sparsity=0.5)
        A = t(X) %*% X
        s = sum(A)
        """
        prog = dml(script).input('$nr', 1e6).output('s')
        s = ml.execute(prog).get('s')
        print s
```

```
[Stage 0:>                                     (0 +
0) / 59]
[Stage 0:>                                     (0 +
1) / 59]
[Stage 0:>                                     (0 +
10) / 59]
[Stage 0:>                                     (1 +
10) / 59]
[Stage 0:=====>                               (11 +
10) / 59]
[Stage 0:=====>                               (18 +
10) / 59]
[Stage 0:=====>                               (25 +
10) / 59]
[Stage 0:=====>                               (34 +
10) / 59]
[Stage 0:=====>                               (44 +
10) / 59]
[Stage 0:=====>                               (54 +
5) / 59]
SystemML Statistics:
Total execution time:          14.134 sec.
Number of executed Spark inst: 2.
```

62608781691.5


Load diabetes dataset from scikit-learn


```
In [8]: %matplotlib inline
```

```
In [9]: diabetes = datasets.load_diabetes()
        diabetes_X = diabetes.data[:, np.newaxis, 2]
        diabetes_X_train = diabetes_X[:-20]
        diabetes_X_test = diabetes_X[-20:]
        diabetes_y_train = np.matrix(diabetes.target[:-20]).T
        diabetes_y_test = np.matrix(diabetes.target[-20:]).T

        plt.scatter(diabetes_X_train, diabetes_y_train, color='black')
        plt.scatter(diabetes_X_test, diabetes_y_test, color='red')
```

```
Out[9]: <matplotlib.collections.PathCollection at 0x7f63be96d850>
```





```
In [10]: diabetes_y_train
```

```
Out[10]: matrix([[ 151.],
 [   75.],
 [ 141.],
 [ 206.],
 [ 135.],
 [   97.],
 [ 138.],
 [   63.],
 [ 110.],
 [ 310.],
 [ 101.],
 [   69.],
 [ 179.],
 [ 185.],
 [ 118.],
 [ 171.],
 [ 166.],
 [ 144.],
 [   97.],
 [ 168.],
 [   68.],
 [   49.],
 [   68.],
 [ 245.],
 [ 184.],
 [ 202.],
 [ 137.],
 [   85.],
 [ 131.],
 [ 283.],
 [ 129.],
 [   59.],
 [ 341.],
 [   87.],
 [   65.],
 [ 102.],
 [ 265.],
 [ 276.],
 [ 252.],
 [   90.],
 [ 100.],
 [   55.],
 [   61.],
 [   92.],
 [ 259.]])
```

[53.],
[190.],
[142.],
[75.],
[142.],
[155.],
[225.],
[59.],
[104.],
[182.],
[128.],
[52.],
[37.],
[170.],
[170.],
[61.],
[144.],
[52.],
[128.],
[71.],
[163.],
[150.],
[97.],
[160.],
[178.],
[48.],
[270.],
[202.],
[111.],
[85.],
[42.],
[170.],
[200.],
[252.],
[113.],
[143.],
[51.],
[52.],
[210.],
[65.],
[141.],
[55.],
[134.],
[42.],
[111.],
[98.],
[164.],
[48.],
[96.],
[90.],
[162.],
[150.],
[279.],
[92.],
[83.],
[128.],
[102.],
[302.],
[198.],
[95.],
[53.],
r 134 1

[144.],
[232.],
[81.],
[104.],
[59.],
[246.],
[297.],
[258.],
[229.],
[275.],
[281.],
[179.],
[200.],
[200.],
[173.],
[180.],
[84.],
[121.],
[161.],
[99.],
[109.],
[115.],
[268.],
[274.],
[158.],
[107.],
[83.],
[103.],
[272.],
[85.],
[280.],
[336.],
[281.],
[118.],
[317.],
[235.],
[60.],
[174.],
[259.],
[178.],
[128.],
[96.],
[126.],
[288.],
[88.],
[292.],
[71.],
[197.],
[186.],
[25.],
[84.],
[96.],
[195.],
[53.],
[217.],
[172.],
[131.],
[214.],
[59.],
[70.],
[220.],

[268.],
[152.],
[47.],
[74.],
[295.],
[101.],
[151.],
[127.],
[237.],
[225.],
[81.],
[151.],
[107.],
[64.],
[138.],
[185.],
[265.],
[101.],
[137.],
[143.],
[141.],
[79.],
[292.],
[178.],
[91.],
[116.],
[86.],
[122.],
[72.],
[129.],
[142.],
[90.],
[158.],
[39.],
[196.],
[222.],
[277.],
[99.],
[196.],
[202.],
[155.],
[77.],
[191.],
[70.],
[73.],
[49.],
[65.],
[263.],
[248.],
[296.],
[214.],
[185.],
[78.],
[93.],
[252.],
[150.],
[77.],
[208.],
[77.],
[108.],
[160.],
[53.].

[55.],
[220.],
[154.],
[259.],
[90.],
[246.],
[124.],
[67.],
[72.],
[257.],
[262.],
[275.],
[177.],
[71.],
[47.],
[187.],
[125.],
[78.],
[51.],
[258.],
[215.],
[303.],
[243.],
[91.],
[150.],
[310.],
[153.],
[346.],
[63.],
[89.],
[50.],
[39.],
[103.],
[308.],
[116.],
[145.],
[74.],
[45.],
[115.],
[264.],
[87.],
[202.],
[127.],
[182.],
[241.],
[66.],
[94.],
[283.],
[64.],
[102.],
[200.],
[265.],
[94.],
[230.],
[181.],
[156.],
[233.],
[60.],
[219.],
[80.],
[68.],
[332.],

[248.],
[84.],
[200.],
[55.],
[85.],
[89.],
[31.],
[129.],
[83.],
[275.],
[65.],
[198.],
[236.],
[253.],
[124.],
[44.],
[172.],
[114.],
[142.],
[109.],
[180.],
[144.],
[163.],
[147.],
[97.],
[220.],
[190.],
[109.],
[191.],
[122.],
[230.],
[242.],
[248.],
[249.],
[192.],
[131.],
[237.],
[78.],
[135.],
[244.],
[199.],
[270.],
[164.],
[72.],
[96.],
[306.],
[91.],
[214.],
[95.],
[216.],
[263.],
[178.],
[113.],
[200.],
[139.],
[139.],
[88.],
[148.],
[88.],
[243.],
[71.],
[77.].

[109.],
[272.],
[60.],
[54.],
[221.],
[90.],
[311.],
[281.],
[182.],
[321.],
[58.],
[262.],
[206.],
[233.],
[242.],
[123.],
[167.],
[63.],
[197.],
[71.],
[168.],
[140.],
[217.],
[121.],
[235.],
[245.],
[40.],
[52.],
[104.],
[132.],
[88.],
[69.],
[219.],
[72.],
[201.],
[110.],
[51.],
[277.],
[63.],
[118.],
[69.],
[273.],
[258.],
[43.],
[198.],
[242.],
[232.],
[175.],
[93.],
[168.],
[275.],
[293.],
[281.],
[72.],
[140.],
[189.],
[181.],
[209.],
[136.],
[261.],
[113.],
-

```
[ 131. ],
[ 174. ],
[ 257. ],
[ 55. ],
[ 84. ],
[ 42. ],
[ 146. ],
[ 212. ]])
```

Example 2: Implement three different algorithms to train linear regression model

Algorithm 1: Linear Regression - Direct Solve (no regularization)

Preliminaries

1. The builtin function `solve(A, b)` computes the least squares solution for system of linear equations

$$Ax = b$$

for the vector x such that

$$\| Ax - b \|$$

is minimized. It is important to note that this function can operate only on small-to-medium sized input matrix that can fit in the driver memory. See the [DML language reference](http://apache.github.io/systemml/dml-language-reference.html) (<http://apache.github.io/systemml/dml-language-reference.html>) for more details.

2. Linear regression model assumes that relationship between input explanatory (feature) variables X and numerical response variable y is linear. The goal is to estimate regression coefficient w (and residual variable) such that

$$y = \text{Normal}(Xw, \sigma^2)$$

$$\text{Cost function, } J(w) = \frac{1}{2}(Xw - y)^2$$

Differentiating with respect to w ,

$$\begin{aligned} dw &= \frac{\partial}{\partial w} \frac{1}{2}(Xw - y)^2 \\ &= \frac{1}{2} 2X^T(Xw - y) \\ &= (X^T X)w - X^T y \end{aligned}$$

Setting the gradient

To find minima, we set the derivative with respect to w to zero,

$$(X^T X)w - (X^T y) = 0$$

$$w = (X^T X)^{-1}(X^T y)$$

$$\text{Let } A = X^T X$$

$$\text{and } b = X^T y$$

$$\text{Therefore, } w = \text{solve}(A, b)$$

```

In [12]: script =
        # add constant feature to X to model intercept
        ones = matrix(1, rows=nrow(X), cols=1)
        X = cbind(X, ones)
        A = t(X) %*% X
        b = t(X) %*% y
        w = solve(A, b)
        bias = as.scalar(w[nrow(w),1])
        w = w[1:nrow(w)-1,]
        """

```

```

In [13]: prog = dml(script).input(X=diabetes_X_train, y=diabetes_y_train).output
        ('w', 'bias')
        w, bias = ml.execute(prog).get('w','bias')
        w = w.toNumPy()

```

SystemML Statistics:
 Total execution time: 0.028 sec.
 Number of executed Spark inst: 2.


```

In [14]: plt.scatter(diabetes_X_train, diabetes_y_train, color='black')
        plt.scatter(diabetes_X_test, diabetes_y_test, color='red')

        plt.plot(diabetes_X_test, (w*diabetes_X_test)+bias, color='blue', lines
        tyle = 'dotted')

```

Out[14]: [<matplotlib.lines.Line2D at 0x7f6348328d10>]




Algorithm 2: Linear Regression - Batch Gradient Descent (no regularization)

Algorithm

```

Step 1: Start with an initial point
while(not converged) {
    Step 2: Compute gradient dw.
    Step 3: Compute stepsize alpha.
    Step 4: Update: w_new = w_old - alpha*dw
}

```



Gradient formula

$$dw = r = (X^T X)w - (X^T y)$$

Step size formula

We perform a line search to choose the step size α to minimize the cost function $J(w)$. From basic calculus, α minimizes the function $J(w)$ when the directional derivative with respect to α is zero.

$$\alpha = \frac{r^T r}{r^T X^T X r}$$

```
In [15]: script = """
# add constant feature to X to model intercepts
ones = matrix(1, rows=nrow(X), cols=1)
X = cbind(X, ones)
max_iter = 100
w = matrix(0, rows=ncol(X), cols=1)
for(i in 1:max_iter){
  XtX = t(X) %*% X
  dw = XtX %*% w - t(X) %*% y
  alpha = (t(dw) %*% dw) / (t(dw) %*% XtX %*% dw)
  w = w - dw*alpha
}
bias = as.scalar(w[nrow(w),1])
w = w[1:nrow(w)-1,]
"""
```


```
In [16]: prog = dml(script).input(X=diabetes_X_train, y=diabetes_y_train).output
('w').output('bias')
w, bias = ml.execute(prog).get('w', 'bias')
w = w.toNumPy()
```

SystemML Statistics:
 Total execution time: 0.081 sec.
 Number of executed Spark inst: 2.

```
In [17]: plt.scatter(diabetes_X_train, diabetes_y_train, color='black')
plt.scatter(diabetes_X_test, diabetes_y_test, color='red')

plt.plot(diabetes_X_test, (w*diabetes_X_test)+bias, color='red', linestyle='dashed')
```

Out[17]: [<matplotlib.lines.Line2D at 0x7f6348296710>]



Algorithm 3: Linear Regression - Conjugate Gradient Gradient (no regularization)

Problem with gradient descent: Takes very similar directions many times

Solution: Enforce conjugacy

Step 1: Start with an initial point

while(not converged) {

 Step 2: Compute gradient ∇w .

 Step 3: Compute stepsize α .

 Step 4: Compute next direction p by enforcing conjugacy with previous direction.

 Step 4: Update: $w_{\text{new}} = w_{\text{old}} + \alpha * p$

}




Figure 6.14: Steepest descent vs. conjugate gradient.

```
In [18]: script = """
# add constant feature to X to model intercepts
X = cbind(X, matrix(1, rows=nrow(X), cols=1))
m = ncol(X); i = 1;
max_iter = 20;
w = matrix(0, rows = m, cols = 1); # initialize weights to 0
dw = - t(X) %*% y; p = - dw;          # dw = (X'X)w - (X'y)
norm_r2 = sum(dw ^ 2);
for(i in 1:max_iter) {
  q = t(X) %*% (X %*% p)
  alpha = norm_r2 / sum(p * q); # Minimizes f(w - alpha*r)
  w = w + alpha * p;           # update weights
  dw = dw + alpha * q;
  old_norm_r2 = norm_r2; norm_r2 = sum(dw ^ 2);
  p = -dw + (norm_r2 / old_norm_r2) * p; # next direction - conjugacy to previous direction
  i = i + 1;
}
bias = as.scalar(w[nrow(w),1])
w = w[1:nrow(w)-1,]
"""
```

```
In [19]: prog = dml(script).input(X=diabetes_X_train, y=diabetes_y_train).output('w').output('bias')
w, bias = ml.execute(prog).get('w','bias')
w = w.toNumPy()
```

SystemML Statistics:


Total execution time: 0.007 sec.

Number of executed Spark inst: 2.

```
In [20]: plt.scatter(diabetes_X_train, diabetes_y_train, color='black')
plt.scatter(diabetes_X_test, diabetes_y_test, color='red')

plt.plot(diabetes_X_test, (w*diabetes_X_test)+bias, color='red', linestyle='dashed')
```

Out[20]: [<matplotlib.lines.Line2D at 0x7f634816a550>]



Example 3: Invoke existing SystemML algorithm

script LinearRegDS.dml using MLContext API


```
In [21]: from systemml import dmlFromResource
prog = dmlFromResource('scripts/algorithms/LinearRegDS.dml').input(X=diabetes_X_train, y=diabetes_y_train).input('$icpt',1.0).output('beta_out')
w = ml.execute(prog).get('beta_out')
w = w.toNumPy()
bias=w[1]
```

```
BEGIN LINEAR REGRESSION SCRIPT
Reading X and Y...
Calling the Direct Solver...
Computing the statistics...
AVG_TOT_Y,153.36255924170615
STDEV_TOT_Y,77.21853383600028
AVG_RES_Y,3.633533705616816E-14
STDEV_RES_Y,63.038506337610244
DISPERSION,3973.853281276927
R2,0.3351312506863875
ADJUSTED_R2,0.33354822985468835
R2_NOBIAS,0.3351312506863875
ADJUSTED_R2_NOBIAS,0.33354822985468835
Writing the output matrix...
END LINEAR REGRESSION SCRIPT
SystemML Statistics:
Total execution time:          0.008 sec.
Number of executed Spark inst: 2.
```

```
In [22]: plt.scatter(diabetes_X_train, diabetes_y_train, color='black')
plt.scatter(diabetes_X_test, diabetes_y_test, color='red')

plt.plot(diabetes_X_test, (w[0]*diabetes_X_test)+bias, color='red', linestyle='dashed')
```

Out[22]: [<matplotlib.lines.Line2D at 0x7f63480b9810>]



Example 4: Invoke existing SystemML algorithm using scikit-learn/SparkML pipeline like API

mllearn API allows a Python programmer to invoke SystemML's algorithms using scikit-learn like API as well as Spark's MLPipeline API.

```
In [23]: from pyspark.sql import SQLContext
        from systemml.mllearn import LinearRegression
        sqlCtx = SQLContext(sc)
```

```
In [24]: regr = LinearRegression(sqlCtx)
        # Train the model using the training sets
        regr.fit(diabetes_X_train, diabetes_y_train)

BEGIN LINEAR REGRESSION SCRIPT
Reading X and Y...
Running the CG algorithm...
||r|| initial value = 64725.64237405237, target value = 0.064725642374
05237
Iteration 1: ||r|| / ||r init|| = 0.013822097283108787
Iteration 2: ||r|| / ||r init|| = 5.369915930350396E-14
The CG algorithm is done.
Computing the statistics...
AVG_TOT_Y,153.36255924170615
STDEV_TOT_Y,77.21853383600028
AVG_RES_Y,-8.227243004822623E-12
STDEV_RES_Y,63.03850633759284
DISPERSION,3973.853281274733
R2,0.33513125068675453
ADJUSTED_R2,0.3335482298550564
R2_NOBIAS,0.33513125068675453
ADJUSTED_R2_NOBIAS,0.3335482298550564
Writing the output matrix...
END LINEAR REGRESSION SCRIPT
SystemML Statistics:
Total execution time:          0.003 sec.
Number of executed Spark inst: 2.
```

Out[24]: lr


```
In [25]: predictions = regr.predict(diabetes_X_test)
```


```
SystemML Statistics:
Total execution time:          0.000 sec.
Number of executed Spark inst: 1.
```

```
In [26]: # Use the trained model to perform prediction
        %matplotlib inline
        plt.scatter(diabetes_X_train, diabetes_y_train, color='black')
        plt.scatter(diabetes_X_test, diabetes_y_test, color='red')

        plt.plot(diabetes_X_test, predictions, color='black')
```

Out[26]: [<matplotlib.lines.Line2D at 0x7f632e6e5b90>]





(Optional) Install OpenBLAS

`!wget https://github.com/xianyi/OpenBLAS/archive/v0.2.20.tar.gz !tar -xzf v0.2.20.tar.gz !cd OpenBLAS-0.2.20/ && make clean !cd OpenBLAS-0.2.20/ && make USE_OPENMP=1`

Example 5: Invoking a Keras model with SystemML

See [SystemML's deep learning documentation \(http://apache.github.io/systemml/deep-learning\)](http://apache.github.io/systemml/deep-learning) for more detail.

```
In [ ]: from mlxtend.data import mnist_data
import numpy as np
from sklearn.utils import shuffle
# Download the MNIST dataset
X, y = mnist_data()
X, y = shuffle(X, y)
# Split the data into training and test
n_samples = len(X)
X_train = X[:int(.9 * n_samples)]
y_train = y[:int(.9 * n_samples)]
X_test = X[int(.9 * n_samples):]
y_test = y[int(.9 * n_samples):]
from keras.models import Sequential
from keras.layers import Input, Dense, Conv2D, MaxPooling2D, Dropout, Flatten
from keras import backend as K
from keras.models import Model
input_shape = (1,28,28) if K.image_data_format() == 'channels_first' else (28,28, 1)
keras_model = Sequential()
keras_model.add(Conv2D(32, kernel_size=(5, 5), activation='relu', input_shape=input_shape, padding='same'))
keras_model.add(MaxPooling2D(pool_size=(2, 2)))
keras_model.add(Conv2D(64, (5, 5), activation='relu', padding='same'))
keras_model.add(MaxPooling2D(pool_size=(2, 2)))
keras_model.add(Flatten())
keras_model.add(Dense(512, activation='relu'))
keras_model.add(Dropout(0.5))
keras_model.add(Dense(10, activation='softmax'))

# Scale the input features
scale = 0.00390625
X_train = X_train*scale
```

```

X_train = X_train*scale
X_test = X_test*scale

from systemml.mllearn import Keras2DML
sysml_model = Keras2DML(spark, keras_model, input_shape=(1,28,28), weights='weights_dir')
sysml_model.setConfigProperty('sysml.native.blas', 'openblas')
sysml_model.setConfigProperty('sysml.native.blas.directory', os.path.join(os.getcwd(), 'OpenBLAS-0.2.20/'))
# sysml_model.setGPU(True).setForceGPU(True)
sysml_model.summary()
sysml_model.fit(X_train, y_train)

```

Using TensorFlow backend.

Loading the model from weights_dir...

SystemML Statistics:

Total execution time: 0.000 sec.

Number of executed Spark inst: 0.

```

[Stage 9:=====> (2
+ 1) / 3]
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|          Name|          Type|          Output|          Weight|          B
ias|          Top|          Bottom|Memory* (train/test)|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
|   conv2d_1_input|          Data| (, 1, 28, 28)|          |
|conv2d_1_input,label|          |          1/0|
|   conv2d_1| Convolution|(, 32, 28, 28)| [32 X 25]| [32 X
11|   conv2d 1|   conv2d 1 input|          25/13|

```