

Dismiss

Join GitHub today

GitHub is home to over 28 million developers working together to host and review code, manage projects, and build software together.


Sign up

Branch: master

[coursera](#) / [coursera\\_ai](#) / [week2](#) / [pytorch](#) / [pytorch\\_tutorial.ipynb](#)

Find file

Copy path

 **Romeo Kienzler** merge from romeokienzler repo

53f3fa0 on Jun 21

[0 contributors](#)

1 lines (1 sloc)530 KB

```
In [5]: %matplotlib inline
```

# Introduction to PyTorch

## PyTorch's tensor library

The most of PyTorch operations are running on **tensors**. A tensor is an multidimensional array. Lets have a look on some basic tensor operations. But first, lets import some important PyTorch libraries:

- **torch** - a Tensor library similar to NumPy, with strong GPU support
- **torch.autograd** - a "tape-based" (about this - later on) automatic differentiation library
- **torch.nn** - a neural networks library deeply integrated with autograd
- **torch.optim** - an optimization package to be used with torch.nn with standard optimization methods such as SGD, RMSProp, LBFGS, Adam etc.

We also set a seed to be able to reproduce the same results later.

```
In [44]: import torch
import torch.autograd as autograd
import torch.nn as nn
import torch.optim as optim

torch.manual_seed(123)

Out[44]: <torch._C.Generator at 0x7fe41fd172f0>
```

## Creating Tensors

Tensors can be created from Python lists with the **torch.Tensor()** function.

```
In [47]: # Create a torch.Tensor object from python list
v = [1, 2, 3]
print(type(v))
v_tensor = torch.Tensor(v)
print(v_tensor)

# Create a torch.Tensor object of size 2x3 from 2x3 matrix
m2x3 = [[1, 2, 3], [4, 5, 6]]
m2x3_tensor = torch.Tensor(m2x3)
print(m2x3_tensor)

# Create a 3D torch.Tensor object of size 3x3x3.
m3x3x3 = [[[1, 2, 3], [4, 5, 6], [7, 8, 9]],
           [[10, 11, 12], [13, 14, 15], [16, 17, 18]],
           [[19, 20, 21], [22, 23, 24], [25, 26, 27]]]
m3x3x3_tensor = torch.Tensor(m3x3x3)
print(m3x3x3_tensor)

#Create a 4D tensor from random data and given dimensions (in this case 3x4x5x6) with torch.randn()
m4x3x3x3_tensor = torch.randn((4, 3, 3, 3))
m4x3x3x3_tensor.shape
print(m4x3x3x3_tensor)

<type 'list'>

1
2
3
[torch.FloatTensor of size 3]

1  2  3
4  5  6
[torch.FloatTensor of size 2x3]

(0 ,.,.) =
1  2  3
4  5  6
7  8  9

(1 ,.,.) =
10 11 12
13 14 15
16 17 18

(2 ,.,.) =
19 20 21
22 23 24
```

```

  25  26  27
[torch.FloatTensor of size 3x3x3]

(0 ,0 ,.,.) =
  0.4934 -0.2766  0.2439
 -1.2116 -0.1520  0.1509
 -0.6251 -0.4416  0.3208

(0 ,1 ,.,.) =
 -0.3273 -0.5305 -0.0172
  0.4719  0.5671  2.7930
  0.3229  0.8552  0.7492

(0 ,2 ,.,.) =
 -1.7119  0.6025 -0.7018
 -1.3130  0.1574  2.0114
  0.1004  0.8222 -0.0176

(1 ,0 ,.,.) =
  1.2481 -0.0710  2.1627
  1.5215 -1.0547  1.7822
  1.9736 -0.3101 -0.8211

(1 ,1 ,.,.) =
  0.1315 -0.6948 -0.5823
  1.0035 -1.4613  0.8985
  0.6210 -0.9679  0.6740

(1 ,2 ,.,.) =
 -1.2828 -0.5097  0.1464
 -0.4860 -0.7529  1.6989
  0.4991 -2.1702  0.5130

(2 ,0 ,.,.) =
 -1.9029  0.8260 -0.6644
  1.6663 -0.5704  1.3906
 -1.4855  0.2987 -0.3029

(2 ,1 ,.,.) =
  0.3354  1.0599  0.1941
 -0.9295  0.5329 -1.1307
  0.1024  2.5200 -1.2324

(2 ,2 ,.,.) =
 -0.8294  0.4342 -0.7374
  0.1591 -1.3560  0.5513
  0.3732  1.4246 -0.6860

(3 ,0 ,.,.) =
 -0.4791  2.2610  1.3191
  1.0130  1.6650  0.8469
  0.2564  0.8049 -0.3107

(3 ,1 ,.,.) =
 -0.9208 -0.1507 -0.6977
  1.5227  2.4308  1.5695
 -0.0657 -0.4277  0.2537

(3 ,2 ,.,.) =
  0.7337  1.1388  0.5516
 -0.6187 -0.6503  0.3017
  0.6113  0.4344 -0.3423
[torch.FloatTensor of size 4x3x3x3]

```

## What is a multidimensional tensor?

Since we frequently deal with  $n > 3$  dimensional tensors, its understanding is very important. The best way to think of a higher ( $n$ ) dimensional object (and tensor in particular) is as of a container which keeps a series of  $n-1$  dimensional objects "inside" of it. We can "pull out" these "inner" objects by indexing into to higher dimensional tensor container. Let's have a look on some examples:

- For a vector  $v$  ( $\dim(v)=1$ ), indexing into it ("pulling out of it") returns its "slice" - a scalar  $s$  ( $\dim(s)=0$ ).
- For a matrix, indexing into it returns its "slice" - a (row or column) vector.
- 3D tensor can be seen as a cube or 3D rectangular consisting of horizontally "stacked" matrices. So if we index into a such tensor it will give us its slice which is a matrix!
- We can't easily visualize 5D (or  $n$ -D) tensors, but the idea is actually the same. If we index in to them, we will pull out an object of dimension  $n-1$ .
- E.g. a 4D tensor can be seen as a list of cubes or 3D reactangulars. If we index in to a 4D tensor, we will get 3D rectangulars.

```
In [46]: # Index into v_tensor and get a scalar
print(v_tensor[0])

# Index into m2x3_tensor and get a vector
print(m2x3_tensor[0])

# Index into m3x3x3_tensor and get a matrix
print(m3x3x3_tensor[0])

# Index into m4x3x3x3_tensor and get a 3D rectangular of size 4x5x6
print(m4x3x3x3_tensor[0])

1.0

1
2
3
[torch.FloatTensor of size 3]

1 2 3
4 5 6
7 8 9
[torch.FloatTensor of size 3x3]

(0 ,.,.) =
-0.1115  0.1204 -0.3696
-0.2404 -1.1969  0.2093
-0.9724 -0.7550  0.3239

(1 ,.,.) =
-0.1085  0.2103 -0.3908
 0.2350  0.6653  0.3528
 0.9728 -0.0386 -0.8861

(2 ,.,.) =
-0.4709 -0.4269 -0.0283
 1.4220 -0.3886 -0.8903
-0.9601 -0.4087  1.0764
[torch.FloatTensor of size 3x3x3]
```

## Operations with Tensors

You can operate on tensors in the ways you would expect. See the documentation <http://pytorch.org/docs/torch.html> (<http://pytorch.org/docs/torch.html>) for a complete list of operations.

Simple mathematical operations: **Addition, Multiplication**

```
In [76]: x = torch.Tensor([1, 2, 3])
y = torch.Tensor([4, 5, 6])
print(x)
print(y)

w = torch.matmul(x, y)
print(w)

1
2
3
[torch.FloatTensor of size 3]

4
5
6
[torch.FloatTensor of size 3]

32.0
```

Helpful operation: **Concatenation**

```
In [78]: # By default, it concatenates along the axis with 0 (rows). It's "stacking" the rows.

x_1 = torch.randn(2, 5)
print(x_1)
y_1 = torch.randn(3, 5)
print(y_1)
z_1 = torch.cat([x_1, y_1])
print(z_1)

# Concatenate columns:
x_2 = torch.randn(2, 3)
```

```

x_2 = torch.randn(2, 5)
print(x_2)
y_2 = torch.randn(3, 5)
print(y_2)
# second arg specifies which axis to concat along. Here we select 1 (columns). It's attaching
the columns.
z_2 = torch.cat([x_2, y_2], 1)
print(z_2)

# If your tensors are not compatible, torch will complain. Uncomment to see the error
torch.cat([x_1, x_2])

0.5374 -0.0899 -0.7969 -0.7968 -0.2611
0.4808 -0.1458 -0.5357 -0.7330 2.3465
[torch.FloatTensor of size 2x5]

-0.4654 1.4965 2.7607 -1.4606 1.0825
-0.9363 -0.2594 -1.3465 0.5388 0.0382
-0.1434 1.4049 0.0889 0.5171 -0.0214
[torch.FloatTensor of size 3x5]

0.5374 -0.0899 -0.7969 -0.7968 -0.2611
0.4808 -0.1458 -0.5357 -0.7330 2.3465
-0.4654 1.4965 2.7607 -1.4606 1.0825
-0.9363 -0.2594 -1.3465 0.5388 0.0382
-0.1434 1.4049 0.0889 0.5171 -0.0214
[torch.FloatTensor of size 5x5]

-0.2001 -0.6503 -0.2019
0.1262 -0.0708 1.4592
[torch.FloatTensor of size 2x3]

0.5823 0.0048 -0.2069 -0.5939 -1.3076
0.9156 0.3748 -1.3108 -2.6207 -1.2696
0.9036 -1.3043 0.3300 -0.6928 -1.5125
[torch.FloatTensor of size 3x5]

RuntimeErrorTraceback (most recent call last)
<ipython-input-78-c85240a5062e> in <module>()
    14 print(y_2)
    15 # second arg specifies which axis to concat along. Here we select 1 (columns). It's at
taching the columns.
--> 16 z_2 = torch.cat([x_2, y_2], 1)
    17 print(z_2)
    18

RuntimeError: inconsistent tensor sizes at /pytorch/torch/lib/TH/generic/THTensorMath.c:2864

```

## Reshaping Tensors

We can use the `.view()` method to reshape a tensor. Often we will need to reshape our data before passing it to a neuronal network.

Let's assume we have 64000 RGB images with the size of 28x28 pixels. We can define an array of shape (64000, 3, 28, 28) to hold them, where 3 is number of color channels:

```

In [97]: x = torch.randn(64000, 3, 28, 28)
# Now we want to add a batch dimension of size 32. We can then infer the second dimension by
placing -1:
x_rehsaped = x.view(32, -1, 3, 28, 28)
print(x_rehsaped.shape)

torch.Size([32, 2000, 3, 28, 28])

```

## Computation Graphs and Automatic Differentiation

A computation graph is a specification of what parameters with which operations are involved in the computation to give the output.

The fundamental class of Pytorch autograd.Variable keeps track of how it was created.

```

In [100]: # Variables wrap tensor objects
x = autograd.Variable(torch.Tensor([1, 2, 3]), requires_grad=True)
# You can access the data with the .data attribute
print(x.data)

y = autograd.Variable(torch.Tensor([4, 5, 6]), requires_grad=True)

# With autograd.Variable you can also perform all the same operations you did with tensors
z = x + y

```

```

z = x + y
print(z.data)

# w knows also that it's result of addition of z elements (AddBackward)
operation = z.grad_fn
print(operation)

1
2
3
[torch.FloatTensor of size 3]

5
7
9
[torch.FloatTensor of size 3]

<AddBackward1 object at 0x7fe41361ffd0>

```

The autograd.Variable knows which operation has created it. But how does that help **compute a gradient**?

```

In [101]: # Lets sum up all the entries in z
s = z.sum()
print(s)
print(s.grad_fn)

Variable containing:
21
[torch.FloatTensor of size 1]

<SumBackward0 object at 0x7fe413507cd0>

```

## Gradient

So now, what is the derivative of this sum with respect to the first component of  $x$ ? Remember, that  $x$  is a tensor of 3 elements:  
 $x = (x_0, x_1, x_2)$

In math, we want a partial derivative of  $s$  with respect to  $x_0$ :  $\frac{\partial s}{\partial x_0}$

Well,  $s$  knows that it was created as a *sum* of the tensor  $z$  elements ( $z_0, z_1, z_2$ ).  $z$  knows that it was the sum  $x + y$ . So

$$s = x_0 + y_0 + x_1 + y_1 + x_2 + y_2 \quad (1)$$

And so  $s$  contains enough information to determine that the derivative of  $s$  with respect to  $x_0$  is 1!

*Reminder:* If you compute the partial derivative with respect to one variable, you handle all other variables as constants. Therefore they all ( $x_1, x_2, y_0, y_1, y_2$ ) get zeroes, and the derivative of  $f(x_0) = x_0$  is 1.

First we need to run **backpropagation** and calculate gradients with respect to every variable. *Note:* if you run backward multiple times, the gradient will increment. That is because Pytorch *accumulates* the gradient into the **.grad property**, since for many models this is very convenient. Lets now have Pytorch compute the gradient, and see that we were right with our guess of 1:

```

In [106]: # calling .backward() on any variable will run backprop, starting from it.
s.backward(retain_graph=True)

```

```

In [107]: print(x)
print(x.grad)
print(y.grad)

Variable containing:
1
2
3
[torch.FloatTensor of size 3]

Variable containing:
3
3
3
[torch.FloatTensor of size 3]

Variable containing:
3
3
3
[torch.FloatTensor of size 3]

```

## How NOT to break the computational graph

Let's create two torch tensors and add them up:

```
In [109]: x = torch.randn((2, 2))
          y = torch.randn((2, 2))
          z = x + y # These are Tensor types, and backprop would not be possible

          print(z)

          0.1730  1.8913
          0.1251  0.1286
          [torch.FloatTensor of size 2x2]
```

Now we wrap the torch tensors in `autograd.Variable`. The `var_z` contains the information for backpropagation:

```
In [111]: var_x = autograd.Variable(x, requires_grad=True)
          var_y = autograd.Variable(y, requires_grad=True)
          # var_z contains enough information to compute gradients, as we saw above
          var_z = var_x + var_y
          print(var_z.grad_fn)

          <AddBackward1 object at 0x7fe41352c110>
```

But what happens if we extract the wrapped tensor object out of `var_z` and re-wrap the tensor in a new `autograd.Variable`?

```
In [112]: var_z_data = var_z.data
          new_var_z = autograd.Variable(var_z_data)
          print(new_var_z.grad_fn)

          None
```

The variable chain is not existing anymore, since we have extracted only data and the whole operations chain was lost. If we try now to compute backward on `new_var_z`, it will throw an error:

```
In [113]: new_var_z.backward(retain_graph=True)

RuntimeErrorTraceback (most recent call last)
<ipython-input-113-9385e2013798> in <module>()
----> 1 new_var_z.backward(retain_graph=True)

/gpfs/fs01/user/s5b4-5d8779714d1900-f211816cda7e/.local/lib/python2.7/site-packages/torch/autograd/variable.pyc in backward(self, gradient, retain_graph, create_graph, retain_variables)
    165         Variable.
    166         """
--> 167         torch.autograd.backward(self, gradient, retain_graph, create_graph, retain_variables)
    168
    169     def register_hook(self, hook):

/gpfs/fs01/user/s5b4-5d8779714d1900-f211816cda7e/.local/lib/python2.7/site-packages/torch/autograd/__init__.pyc in backward(variables, grad_variables, retain_graph, create_graph, retain_variables)
    97
    98     Variable._execution_engine.run_backward(
--> 99         variables, grad_variables, retain_graph)
    100
    101

RuntimeError: element 0 of variables does not require grad and does not have a grad_fn
```

## CUDA

Check whether GPU acceleration with **CUDA** is available

```
In [42]: # Let us run this cell only if CUDA is available
          if torch.cuda.is_available():
              # creates a LongTensor and transfers it
              # to GPU as torch.cuda.LongTensor
              a = torch.LongTensor(10).fill_(3).cuda()
              print(type(a))
              b = a.cpu()
              # transfers it to CPU, back to
              # being a torch.LongTensor
```

## Linear Model

```
In [1]: import torch
          import torch.nn as nn
```

```
from torch.autograd import Variable
import numpy as np
```

```
In [2]: x = [i for i in range(20)] #List comprehension
x_train = np.array(x, dtype=np.float32)
x_train = x_train.reshape(-1, 1)
print(x)
print(x_train.shape)

y = [(5*i + 2) for i in x] #List comprehension
y_train = np.array(y, dtype=np.float32)
y_train = y_train.reshape(-1, 1)
print(y)
print(y_train.shape)

[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19]
(20, 1)
[2, 7, 12, 17, 22, 27, 32, 37, 42, 47, 52, 57, 62, 67, 72, 77, 82, 87, 92, 97]
(20, 1)
```

## Create Model Class

```
In [6]: class LinearRegressor(nn.Module):
        def __init__(self, input_dim, output_dim):
            super(LinearRegressor, self).__init__()
            self.linear = nn.Linear(input_dim, output_dim)

        def forward(self, x):
            out = self.linear(x)
            return out

input_dim = 1
output_dim = 1

model = LinearRegressor(input_dim, output_dim)

model
```

```
Out[6]: LinearRegressor(
  (linear): Linear(in_features=1, out_features=1)
)
```

## Loss & Optimizer

```
In [7]: loss_function = nn.MSELoss()

optimizer = torch.optim.SGD(model.parameters(), lr=0.001)
optimizer
loss_function
```

```
Out[7]: MSELoss(
)
```

```
In [8]: epochs = 500

for epoch in range(epochs):
    epoch += 1
    #Convert inputs and outputs to torch variable
    inputs = Variable(torch.from_numpy(x_train))

    real_outputs = Variable(torch.from_numpy(y_train))

    # Reset Gradients
    optimizer.zero_grad()

    # Forward - compute the output
    pred_outputs = model(inputs)

    # Loss
    loss = loss_function(pred_outputs, real_outputs)

    # Backward - compute gradients
    loss.backward()

    # Update parameters
    optimizer.step()

    print('epoch {}, loss {}'.format(epoch, loss.data[0]))
```

```
epoch 1, loss 3330.96435547
epoch 2, loss 1881.8626709
epoch 3, loss 1063.4017334
epoch 4, loss 601.130065918
```



epoch 5, loss 340.03616333  
epoch 6, loss 192.568328857  
epoch 7, loss 109.277175903  
epoch 8, loss 62.2333488464  
epoch 9, loss 35.6621856689  
epoch 10, loss 20.6541633606  
epoch 11, loss 12.1769609451  
epoch 12, loss 7.38844966888  
epoch 13, loss 4.68333339691  
epoch 14, loss 3.15493202209  
epoch 15, loss 2.29113745689  
epoch 16, loss 1.80271029472  
epoch 17, loss 1.52630400658  
epoch 18, loss 1.36964297295  
epoch 19, loss 1.28061759472  
epoch 20, loss 1.22979319096  
epoch 21, loss 1.20054602623  
epoch 22, loss 1.18348562717  
epoch 23, loss 1.17330884933  
epoch 24, loss 1.16701996326  
epoch 25, loss 1.16292822361  
epoch 26, loss 1.16007828712  
epoch 27, loss 1.15793061256  
epoch 28, loss 1.15617823601  
epoch 29, loss 1.15465211868  
epoch 30, loss 1.15325260162  
epoch 31, loss 1.15192627907  
epoch 32, loss 1.15064108372  
epoch 33, loss 1.14937984943  
epoch 34, loss 1.14813303947  
epoch 35, loss 1.14689469337  
epoch 36, loss 1.14566218853  
epoch 37, loss 1.14443290234  
epoch 38, loss 1.14320623875  
epoch 39, loss 1.14198327065  
epoch 40, loss 1.14075922966  
epoch 41, loss 1.13953721523  
epoch 42, loss 1.13831853867  
epoch 43, loss 1.13709974289  
epoch 44, loss 1.13588225842  
epoch 45, loss 1.13466560841  
epoch 46, loss 1.13345181942  
epoch 47, loss 1.13223946095  
epoch 48, loss 1.13102686405  
epoch 49, loss 1.12981712818  
epoch 50, loss 1.12860763073  
epoch 51, loss 1.12740015984  
epoch 52, loss 1.12619328499  
epoch 53, loss 1.12498831749  
epoch 54, loss 1.12378358841  
epoch 55, loss 1.12258088589  
epoch 56, loss 1.12137913704  
epoch 57, loss 1.12017941475  
epoch 58, loss 1.11898028851  
epoch 59, loss 1.11778283119  
epoch 60, loss 1.11658644676  
epoch 61, loss 1.11539149284  
epoch 62, loss 1.11419713497  
epoch 63, loss 1.11300492287  
epoch 64, loss 1.11181366444  
epoch 65, loss 1.11062431335  
epoch 66, loss 1.10943520069  
epoch 67, loss 1.10824799538  
epoch 68, loss 1.10706198215  
epoch 69, loss 1.10587668419  
epoch 70, loss 1.10469269753  
epoch 71, loss 1.10351061821  
epoch 72, loss 1.10232949257  
epoch 73, loss 1.10114979744  
epoch 74, loss 1.09997189045  
epoch 75, loss 1.09879422188  
epoch 76, loss 1.09761846066  
epoch 77, loss 1.09644293785  
epoch 78, loss 1.09526956081  
epoch 79, loss 1.09409844875  
epoch 80, loss 1.09292626381  
epoch 81, loss 1.09175646305  
epoch 82, loss 1.09058868885  
epoch 83, loss 1.0894215107  
epoch 84, loss 1.08825469017  
epoch 85, loss 1.08709037304  
epoch 86, loss 1.085927248  
epoch 87, loss 1.0847645998  
epoch 88, loss 1.08360350132  
epoch 89, loss 1.08244431019  
epoch 90, loss 1.08128559589  
epoch 91, loss 1.08012807369

epoch 92, loss 1.07897210121  
epoch 93, loss 1.07781684399  
epoch 94, loss 1.07666349411  
epoch 95, loss 1.07551205158  
epoch 96, loss 1.07435965538  
epoch 97, loss 1.0732101202  
epoch 98, loss 1.07206165791  
epoch 99, loss 1.07091379166  
epoch 100, loss 1.06976830959  
epoch 101, loss 1.06862294674  
epoch 102, loss 1.06747913361  
epoch 103, loss 1.06633603573  
epoch 104, loss 1.0651961565  
epoch 105, loss 1.06405568123  
epoch 106, loss 1.06291675568  
epoch 107, loss 1.06177866459  
epoch 108, loss 1.06064271927  
epoch 109, loss 1.05950784683  
epoch 110, loss 1.05837345123  
epoch 111, loss 1.05724036694  
epoch 112, loss 1.05610907078  
epoch 113, loss 1.0549788475  
epoch 114, loss 1.05385005474  
epoch 115, loss 1.05272185802  
epoch 116, loss 1.05159509182  
epoch 117, loss 1.0504693985  
epoch 118, loss 1.04934501648  
epoch 119, loss 1.04822170734  
epoch 120, loss 1.04710018635  
epoch 121, loss 1.04597961903  
epoch 122, loss 1.04486060143  
epoch 123, loss 1.04374194145  
epoch 124, loss 1.04262447357  
epoch 125, loss 1.04150927067  
epoch 126, loss 1.04039406776  
epoch 127, loss 1.03928053379  
epoch 128, loss 1.03816819191  
epoch 129, loss 1.03705775738  
epoch 130, loss 1.03594756126  
epoch 131, loss 1.03483831882  
epoch 132, loss 1.03373169899  
epoch 133, loss 1.03262424469  
epoch 134, loss 1.03151917458  
epoch 135, loss 1.03041529655  
epoch 136, loss 1.02931237221  
epoch 137, loss 1.02821099758  
epoch 138, loss 1.02711033821  
epoch 139, loss 1.02601122856  
epoch 140, loss 1.02491283417  
epoch 141, loss 1.02381563187  
epoch 142, loss 1.02271962166  
epoch 143, loss 1.0216255188  
epoch 144, loss 1.02053201199  
epoch 145, loss 1.01943993568  
epoch 146, loss 1.01834869385  
epoch 147, loss 1.01725888252  
epoch 148, loss 1.01616990566  
epoch 149, loss 1.0150822401  
epoch 150, loss 1.01399600506  
epoch 151, loss 1.01291131973  
epoch 152, loss 1.01182675362  
epoch 153, loss 1.01074433327  
epoch 154, loss 1.00966215134  
epoch 155, loss 1.00858151913  
epoch 156, loss 1.00750219822  
epoch 157, loss 1.0064227581  
epoch 158, loss 1.00534558296  
epoch 159, loss 1.00427126884  
epoch 160, loss 1.00319552422  
epoch 161, loss 1.00212168694  
epoch 162, loss 1.00104928017  
epoch 163, loss 0.99997740984  
epoch 164, loss 0.998907446861  
epoch 165, loss 0.997838199139  
epoch 166, loss 0.996771931648  
epoch 167, loss 0.995703995228  
epoch 168, loss 0.994637966156  
epoch 169, loss 0.993573665619  
epoch 170, loss 0.992510199547  
epoch 171, loss 0.991447746754  
epoch 172, loss 0.990386784077  
epoch 173, loss 0.989326655865  
epoch 174, loss 0.988267123699  
epoch 175, loss 0.987210392952  
epoch 176, loss 0.986153483391  
epoch 177, loss 0.985097527504  
epoch 178, loss 0.984043121338

epoch 179, loss 0.982990145683  
epoch 180, loss 0.981939017773  
epoch 181, loss 0.980888009071  
epoch 182, loss 0.979837536812  
epoch 183, loss 0.978789806366  
epoch 184, loss 0.977740943432  
epoch 185, loss 0.976695179939  
epoch 186, loss 0.975649952888  
epoch 187, loss 0.974604964256  
epoch 188, loss 0.973562121391  
epoch 189, loss 0.972519397736  
epoch 190, loss 0.971478819847  
epoch 191, loss 0.9704387784  
epoch 192, loss 0.969401478767  
epoch 193, loss 0.968362212181  
epoch 194, loss 0.967327296734  
epoch 195, loss 0.966291606426  
epoch 196, loss 0.965257465839  
epoch 197, loss 0.964224636555  
epoch 198, loss 0.963192462921  
epoch 199, loss 0.962162315845  
epoch 200, loss 0.961131751537  
epoch 201, loss 0.960103809834  
epoch 202, loss 0.959075450897  
epoch 203, loss 0.958049118519  
epoch 204, loss 0.957024395466  
epoch 205, loss 0.95599925518  
epoch 206, loss 0.954976439476  
epoch 207, loss 0.953954398632  
epoch 208, loss 0.952932834625  
epoch 209, loss 0.951912999153  
epoch 210, loss 0.95089495182  
epoch 211, loss 0.949876189232  
epoch 212, loss 0.948859095573  
epoch 213, loss 0.947844982147  
epoch 214, loss 0.946830153465  
epoch 215, loss 0.945816338062  
epoch 216, loss 0.944804489613  
epoch 217, loss 0.943793416023  
epoch 218, loss 0.942782759666  
epoch 219, loss 0.941774070263  
epoch 220, loss 0.940765857697  
epoch 221, loss 0.939758956432  
epoch 222, loss 0.938753008842  
epoch 223, loss 0.937748432159  
epoch 224, loss 0.936745524406  
epoch 225, loss 0.935741782188  
epoch 226, loss 0.934740841389  
epoch 227, loss 0.933740496635  
epoch 228, loss 0.932740986347  
epoch 229, loss 0.931742370129  
epoch 230, loss 0.930745720863  
epoch 231, loss 0.92974960804  
epoch 232, loss 0.928754210472  
epoch 233, loss 0.927759647369  
epoch 234, loss 0.926767170429  
epoch 235, loss 0.925775527954  
epoch 236, loss 0.924784362316  
epoch 237, loss 0.923794746399  
epoch 238, loss 0.922805786133  
epoch 239, loss 0.921818375587  
epoch 240, loss 0.920831859112  
epoch 241, loss 0.919846057892  
epoch 242, loss 0.918861567974  
epoch 243, loss 0.917877852917  
epoch 244, loss 0.91689568758  
epoch 245, loss 0.915914237499  
epoch 246, loss 0.914934277534  
epoch 247, loss 0.913955032825  
epoch 248, loss 0.912976861  
epoch 249, loss 0.91199952364  
epoch 250, loss 0.911023259163  
epoch 251, loss 0.910047829151  
epoch 252, loss 0.909074306488  
epoch 253, loss 0.908101379871  
epoch 254, loss 0.907129406929  
epoch 255, loss 0.906159043312  
epoch 256, loss 0.905188858509  
epoch 257, loss 0.904220223427  
epoch 258, loss 0.903252482414  
epoch 259, loss 0.902285277843  
epoch 260, loss 0.901319384575  
epoch 261, loss 0.900355517864  
epoch 262, loss 0.899390697479  
epoch 263, loss 0.898429691792  
epoch 264, loss 0.897467255592  
epoch 265, loss 0.896506667137

epoch 265, loss 0.895547032356  
epoch 266, loss 0.89458835125  
epoch 267, loss 0.893631160259  
epoch 268, loss 0.892674803734  
epoch 269, loss 0.891718745232  
epoch 270, loss 0.890764534473  
epoch 271, loss 0.889811813831  
epoch 272, loss 0.888859570026  
epoch 273, loss 0.887908160686  
epoch 274, loss 0.886957943439  
epoch 275, loss 0.886007785797  
epoch 276, loss 0.885059535503  
epoch 277, loss 0.884112656116  
epoch 278, loss 0.883166611195  
epoch 279, loss 0.882221400738  
epoch 280, loss 0.881277084351  
epoch 281, loss 0.880333602428  
epoch 282, loss 0.879391491413  
epoch 283, loss 0.8784506917  
epoch 284, loss 0.877509474754  
epoch 285, loss 0.876570820808  
epoch 286, loss 0.875633537769  
epoch 287, loss 0.874695897102  
epoch 288, loss 0.873760223389  
epoch 289, loss 0.87282449007  
epoch 290, loss 0.871890425682  
epoch 291, loss 0.870957255363  
epoch 292, loss 0.870025336742  
epoch 293, loss 0.869093596935  
epoch 294, loss 0.868164181709  
epoch 295, loss 0.867234230042  
epoch 296, loss 0.866306006908  
epoch 297, loss 0.865379035473  
epoch 298, loss 0.864453196526  
epoch 299, loss 0.863526701927  
epoch 300, loss 0.862603843212  
epoch 301, loss 0.861680030823  
epoch 302, loss 0.860757827759  
epoch 303, loss 0.85983645916  
epoch 304, loss 0.858916163445  
epoch 305, loss 0.857997298241  
epoch 306, loss 0.85707873106  
epoch 307, loss 0.856161415577  
epoch 308, loss 0.855246067047  
epoch 309, loss 0.854329943657  
epoch 310, loss 0.853415310383  
epoch 311, loss 0.852502346039  
epoch 312, loss 0.851590156555  
epoch 313, loss 0.850677967072  
epoch 314, loss 0.849767506123  
epoch 315, loss 0.848858356476  
epoch 316, loss 0.847949504852  
epoch 317, loss 0.84704208374  
epoch 318, loss 0.846135795116  
epoch 319, loss 0.845229744911  
epoch 320, loss 0.844325423241  
epoch 321, loss 0.843421339989  
epoch 322, loss 0.842518806458  
epoch 323, loss 0.841616034508  
epoch 324, loss 0.840716183186  
epoch 325, loss 0.839816212654  
epoch 326, loss 0.838917732239  
epoch 327, loss 0.838019549847  
epoch 328, loss 0.837123036385  
epoch 329, loss 0.836227238178  
epoch 330, loss 0.835331439972  
epoch 331, loss 0.834438204765  
epoch 332, loss 0.833545207977  
epoch 333, loss 0.832652568817  
epoch 334, loss 0.831760764122  
epoch 335, loss 0.830871462822  
epoch 336, loss 0.829982161522  
epoch 337, loss 0.829094231129  
epoch 338, loss 0.828205943108  
epoch 339, loss 0.827320098877  
epoch 340, loss 0.826433956623  
epoch 341, loss 0.825549960136  
epoch 342, loss 0.824666500092  
epoch 343, loss 0.823783993721  
epoch 344, loss 0.822902083397  
epoch 345, loss 0.822021186352  
epoch 346, loss 0.821141421795  
epoch 347, loss 0.820262610912  
epoch 348, loss 0.819384694099  
epoch 349, loss 0.818507790565  
epoch 350, loss 0.817631900311  
epoch 351, loss 0.816756844534

epoch 352, loss 0.816756844521  
epoch 353, loss 0.815882503986  
epoch 354, loss 0.815009713173  
epoch 355, loss 0.814137279987  
epoch 356, loss 0.813265919685  
epoch 357, loss 0.812395095825  
epoch 358, loss 0.811525523663  
epoch 359, loss 0.810657143593  
epoch 360, loss 0.809789180756  
epoch 361, loss 0.808922946453  
epoch 362, loss 0.808056950569  
epoch 363, loss 0.807192504406  
epoch 364, loss 0.806328296661  
epoch 365, loss 0.805464625359  
epoch 366, loss 0.804603278637  
epoch 367, loss 0.803742587566  
epoch 368, loss 0.802882015705  
epoch 369, loss 0.802022814751  
epoch 370, loss 0.801163971424  
epoch 371, loss 0.800305664539  
epoch 372, loss 0.799450039864  
epoch 373, loss 0.798594295979  
epoch 374, loss 0.797739744186  
epoch 375, loss 0.796885728836  
epoch 376, loss 0.796033501625  
epoch 377, loss 0.7951810956  
epoch 378, loss 0.794329583645  
epoch 379, loss 0.79347962141  
epoch 380, loss 0.792630374432  
epoch 381, loss 0.79178249836  
epoch 382, loss 0.790934920311  
epoch 383, loss 0.790088713169  
epoch 384, loss 0.789242625237  
epoch 385, loss 0.788398146629  
epoch 386, loss 0.78755402565  
epoch 387, loss 0.786711573601  
epoch 388, loss 0.785869300365  
epoch 389, loss 0.785028576851  
epoch 390, loss 0.78418803215  
epoch 391, loss 0.783348798752  
epoch 392, loss 0.782510399818  
epoch 393, loss 0.781673669815  
epoch 394, loss 0.780836701393  
epoch 395, loss 0.780000209808  
epoch 396, loss 0.779165744781  
epoch 397, loss 0.778331875801  
epoch 398, loss 0.77749890089  
epoch 399, loss 0.776666760445  
epoch 400, loss 0.775835394859  
epoch 401, loss 0.775004982948  
epoch 402, loss 0.77417576313  
epoch 403, loss 0.773346424103  
epoch 404, loss 0.772519230843  
epoch 405, loss 0.771692216396  
epoch 406, loss 0.770866215229  
epoch 407, loss 0.770041763783  
epoch 408, loss 0.769217371941  
epoch 409, loss 0.768394172192  
epoch 410, loss 0.767570853233  
epoch 411, loss 0.766750276089  
epoch 412, loss 0.765929281712  
epoch 413, loss 0.765109777451  
epoch 414, loss 0.764290452003  
epoch 415, loss 0.763472735882  
epoch 416, loss 0.762655556202  
epoch 417, loss 0.761839568615  
epoch 418, loss 0.761023163795  
epoch 419, loss 0.760208964348  
epoch 420, loss 0.759395718575  
epoch 421, loss 0.758583366871  
epoch 422, loss 0.757771253586  
epoch 423, loss 0.756960272789  
epoch 424, loss 0.756150126457  
epoch 425, loss 0.755340695381  
epoch 426, loss 0.754532814026  
epoch 427, loss 0.753725349903  
epoch 428, loss 0.752918124199  
epoch 429, loss 0.752112329006  
epoch 430, loss 0.751306772232  
epoch 431, loss 0.750502943993  
epoch 432, loss 0.749699950218  
epoch 433, loss 0.748897194862  
epoch 434, loss 0.748095393181  
epoch 435, loss 0.747294843197  
epoch 436, loss 0.746495068073  
epoch 437, loss 0.745696365833  
epoch 438, loss 0.744898676872

epoch 439, loss 0.744100928307  
epoch 440, loss 0.74330496788  
epoch 441, loss 0.742509245872  
epoch 442, loss 0.741714060307  
epoch 443, loss 0.740920722485  
epoch 444, loss 0.740127801895  
epoch 445, loss 0.739335119724  
epoch 446, loss 0.738543987274  
epoch 447, loss 0.737753987312  
epoch 448, loss 0.73696398735  
epoch 449, loss 0.736175179482  
epoch 450, loss 0.735387027264  
epoch 451, loss 0.734600186348  
epoch 452, loss 0.73381459713  
epoch 453, loss 0.733029007912  
epoch 454, loss 0.732244610786  
epoch 455, loss 0.731460571289  
epoch 456, loss 0.73067766428  
epoch 457, loss 0.729895353317  
epoch 458, loss 0.729114890099  
epoch 459, loss 0.728334486485  
epoch 460, loss 0.727554738522  
epoch 461, loss 0.72677564621  
epoch 462, loss 0.725997924805  
epoch 463, loss 0.725220680237  
epoch 464, loss 0.72444498539  
epoch 465, loss 0.723669946194  
epoch 466, loss 0.722895383835  
epoch 467, loss 0.722121119499  
epoch 468, loss 0.72134822607  
epoch 469, loss 0.720576047897  
epoch 470, loss 0.719805121422